



Chapter 3. Threads and Concurrency

Abraham Silberschatz, Peter Baer Galvin, Greg Gagne.
(2018). *Operating System Concepts* (10th ed.). Wiley.

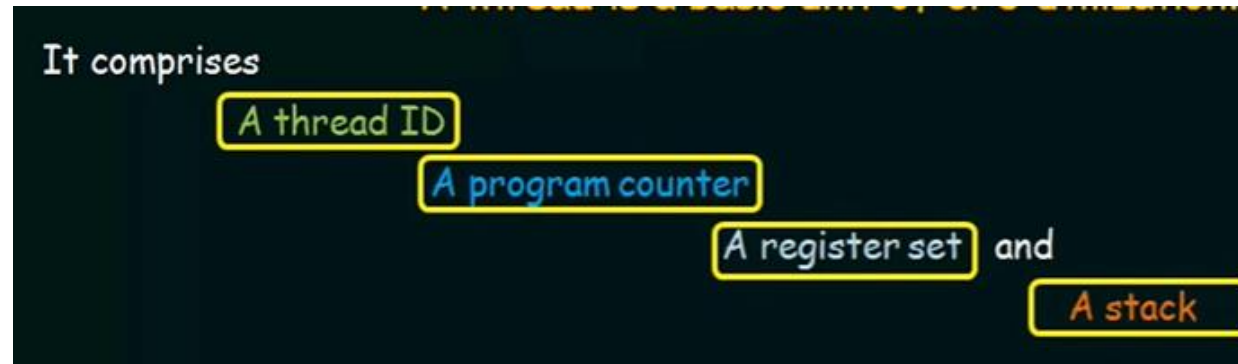


Contents

- 1. Overview
- 2. Multicore Programming
- 3. Multithreading Models
- 4. Thread Libraries

1. Overview

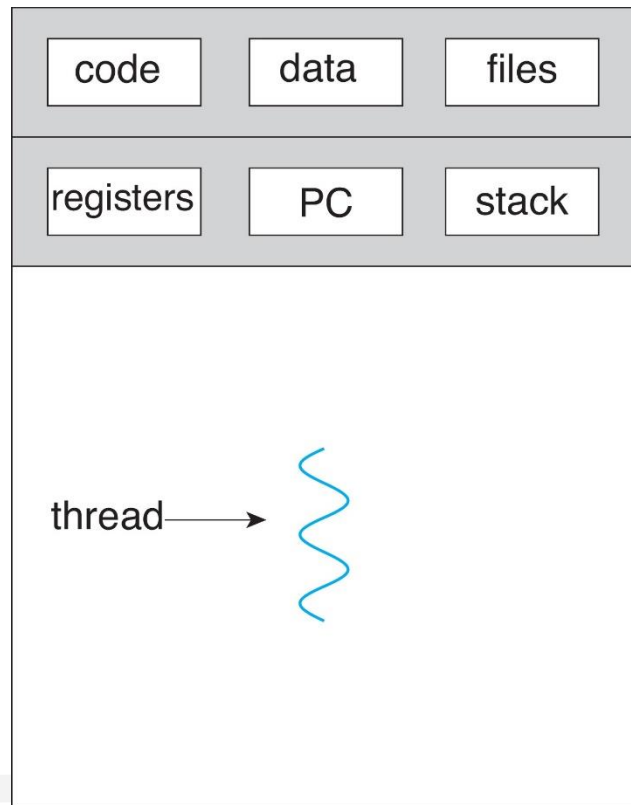
- A thread is a **basic unit of CPU utilization**; it comprises a thread ID, a program counter (PC), a register set, and a stack.



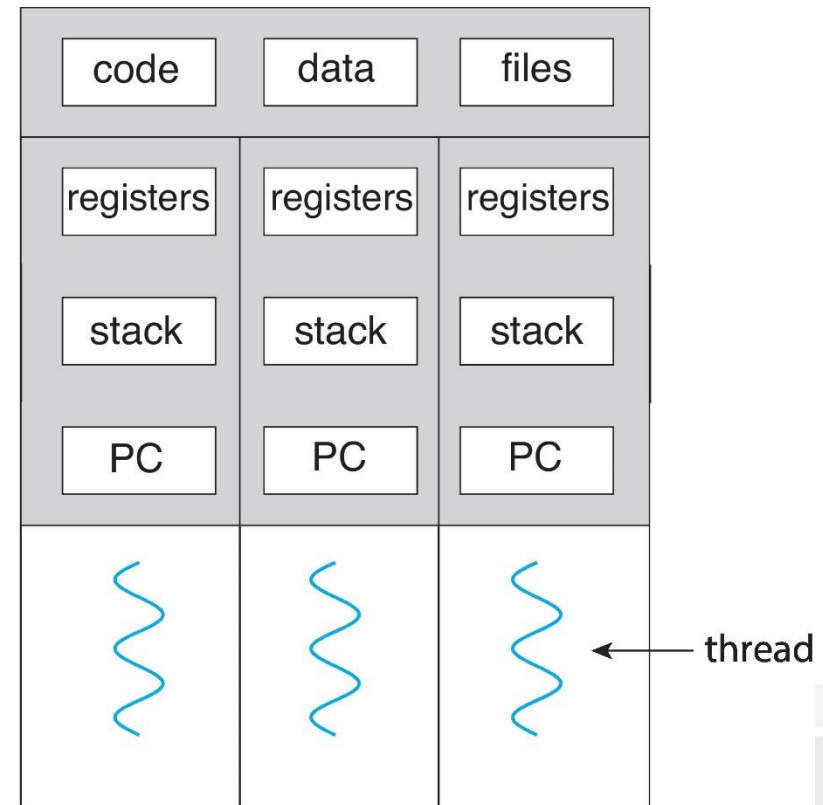
- It shares with other threads belonging to the same process its code section, data section, and other operating-system resources, such as open files and signals.
- A traditional process has a single thread of control. If a process has multiple threads of control, it can perform more than one task at a time.

1. Overview

- Single and multithreaded processes



single-threaded process



multithreaded process

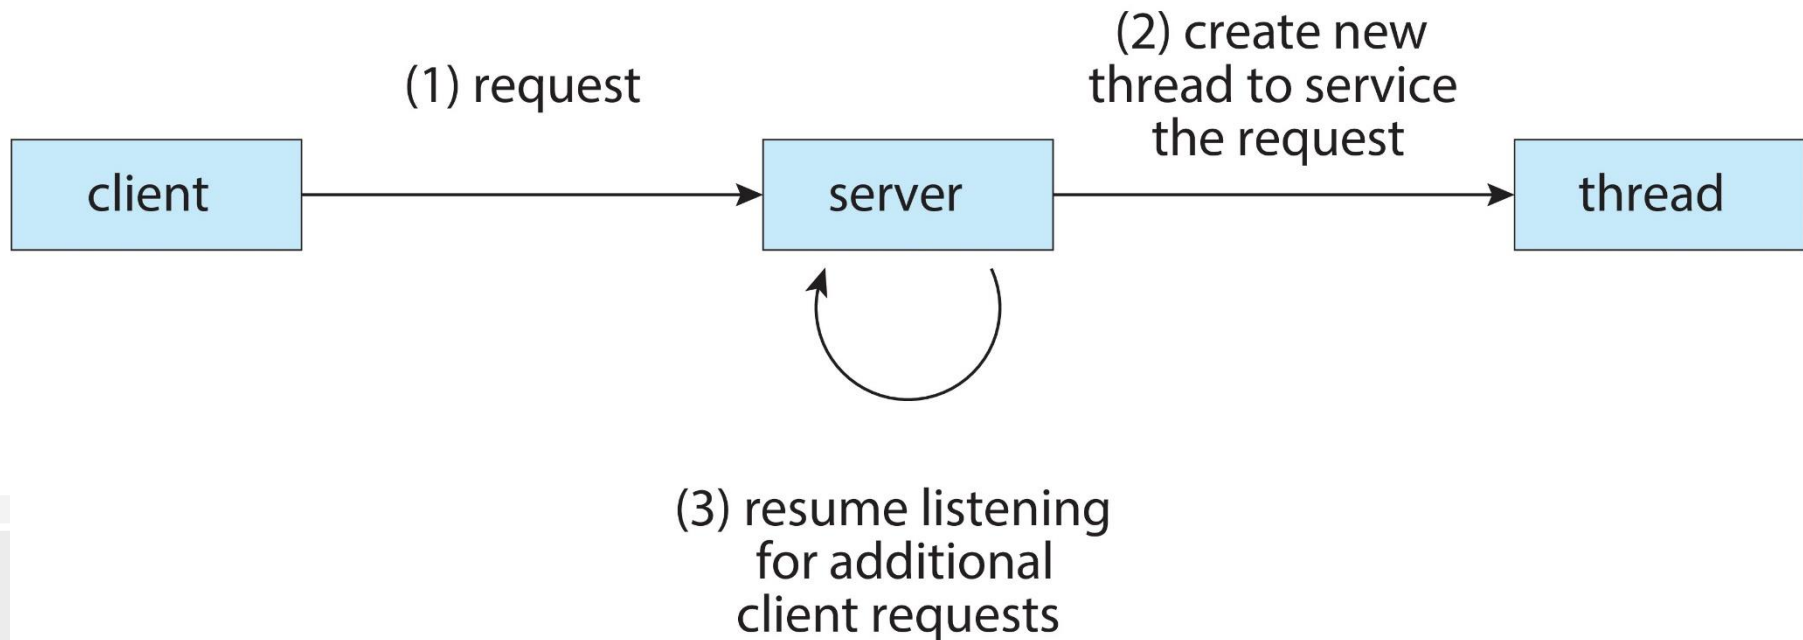


1. Overview

- Most modern applications are multithreaded.
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

1. Overview

- Multithreaded server architecture
 - The server will create a separate thread that listens for client requests.
 - When a request is made, the server creates a new thread to service the request and resumes listening for additional requests.





1. Overview


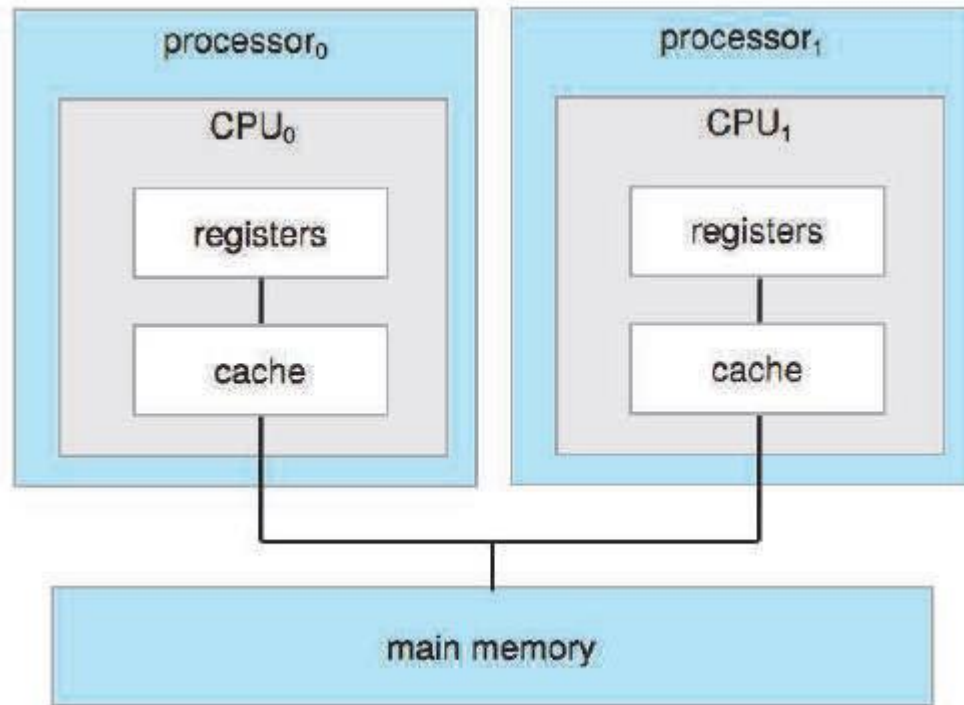
- Benefits
 - **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
 - **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
 - **Economy** – cheaper than process creation, thread switching lower overhead than context switching
 - **Scalability** – process can take advantage of multicore architectures



2. Multicore Programming

- Most modern computers are **multiprocessor systems**.
- The definition of multiprocessor has evolved over time and now includes **multicore systems**, in which multiple computing cores reside on a single chip.
- Multicore systems can be more efficient than multiple chips with single cores because on-chip communication is faster than between-chip communication.

2. Multicore Programming





2. Multicore Programming

- Multicore or multiprocessor systems puts pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously.
- **Concurrency** supports more than one task making progress
 - Single processor/core, scheduler providing concurrency




2. Multicore Programming


- On a system with a single computing core, **concurrency** means that the execution of the threads will be interleaved over time, because the processing core is capable of executing only one thread at a time.
- On a system with multiple cores, however, concurrency means that some threads can run in parallel, because the system can assign a separate thread to each core.
 - These kinds of execution are called Parallelism
- **Parallelism** implies a system can perform more than one task simultaneously.
- **Concurrency** supports more than one task making progress

2. Multicore Programming

- Concurrent execution on single-core system:

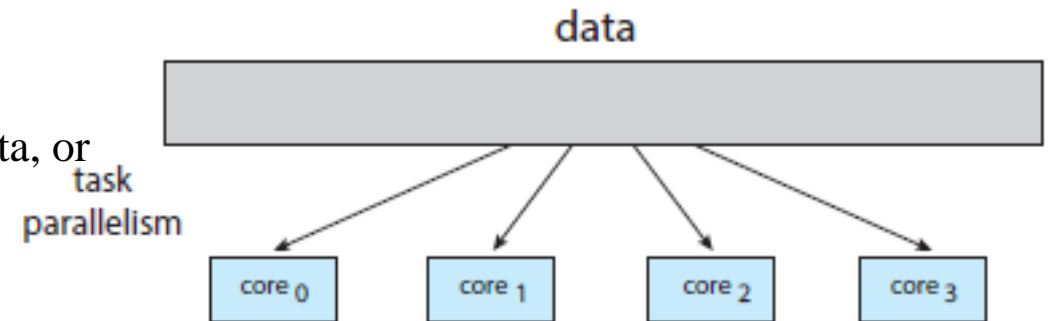
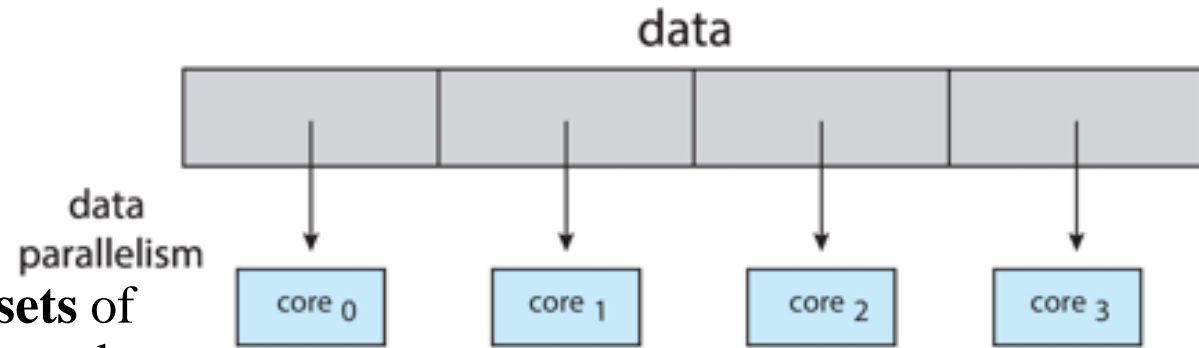


- Parallelism on a multi-core system:



2. Multicore Programming

- Types of parallelism
 - Data parallelism
 - Task parallelism
- **Data parallelism** focuses on distributing **subsets** of the same data across multiple computing cores and performing **the same operation** on each core.
- **Task parallelism** involves distributing not data but tasks (threads) across multiple computing cores.
 - Each thread is performing a unique operation.
 - Different threads may be operating on the same data, or they may be operating on different data





2. Multicore Programming

- Amdahl's Law
 - Amdahl's Law is a formula that identifies potential performance gains from adding additional computing cores to an application that has both serial (nonparallel) and parallel components. If S is the portion of the application that must be performed serially on a system with N processing cores, the formula appears as follows:

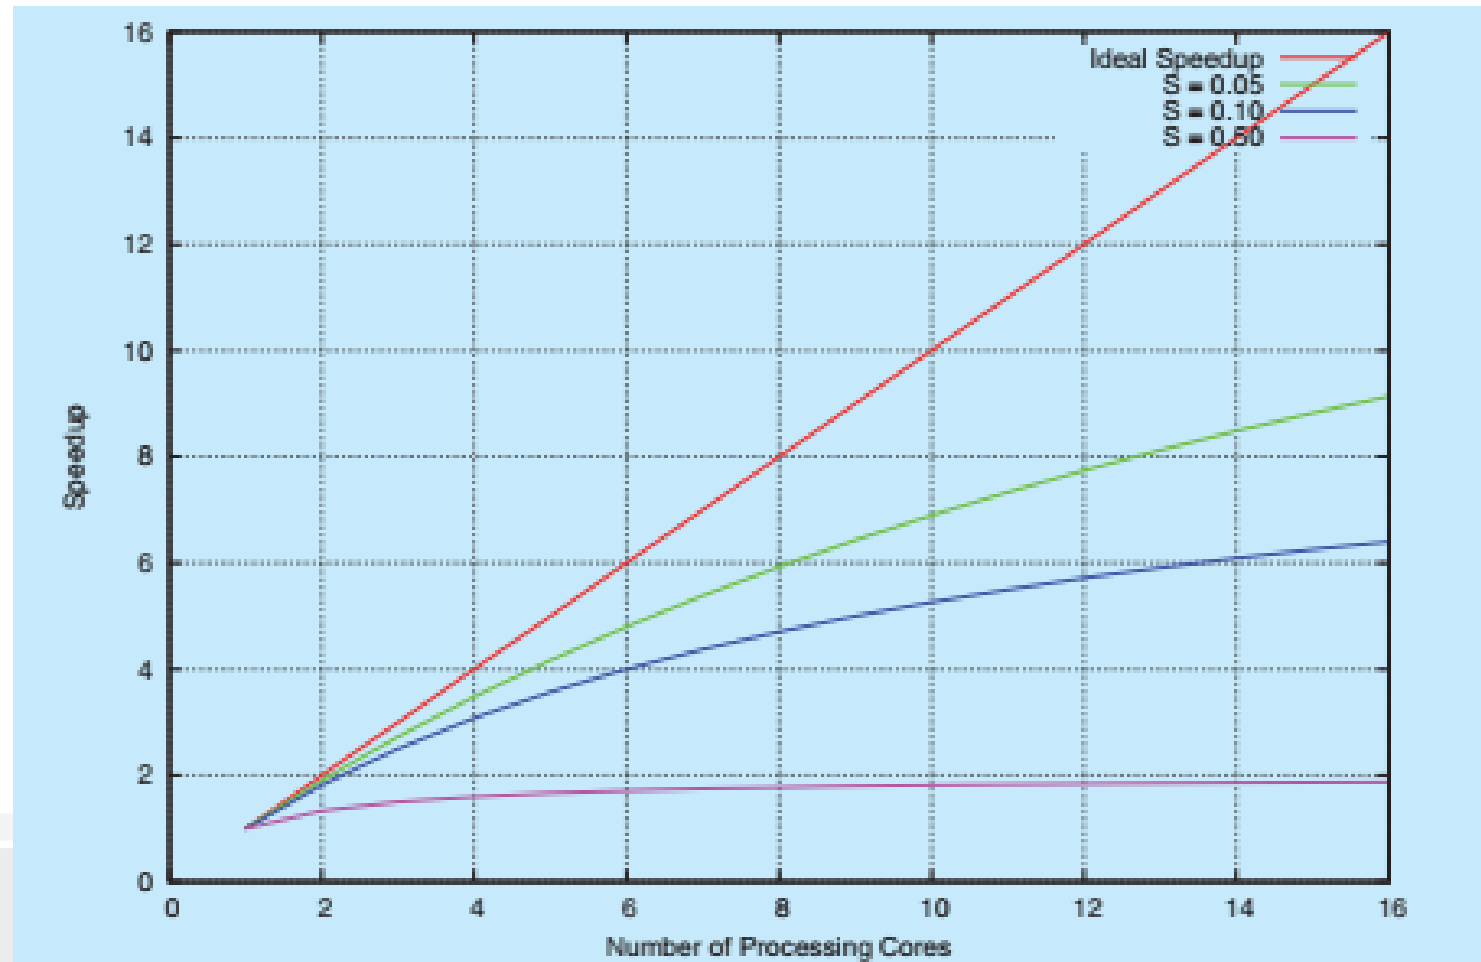
$$speedup \leq \frac{1}{S + \frac{(1 - S)}{N}}$$

- As an example, assume we have an application that is 75 percent parallel and 25 percent serial. If we run this application on a system with two processing cores ($N = 2$), we can get a speedup of 1.6 times. If we add two additional cores ($N = 4$), the speedup is 2.28 times.
- As N approaches infinity, speedup approaches $1 / S$.
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores.

2. Multicore Programming


- Amdahl's Law

$$speedup \leq \frac{1}{S + \frac{(1 - S)}{N}}$$

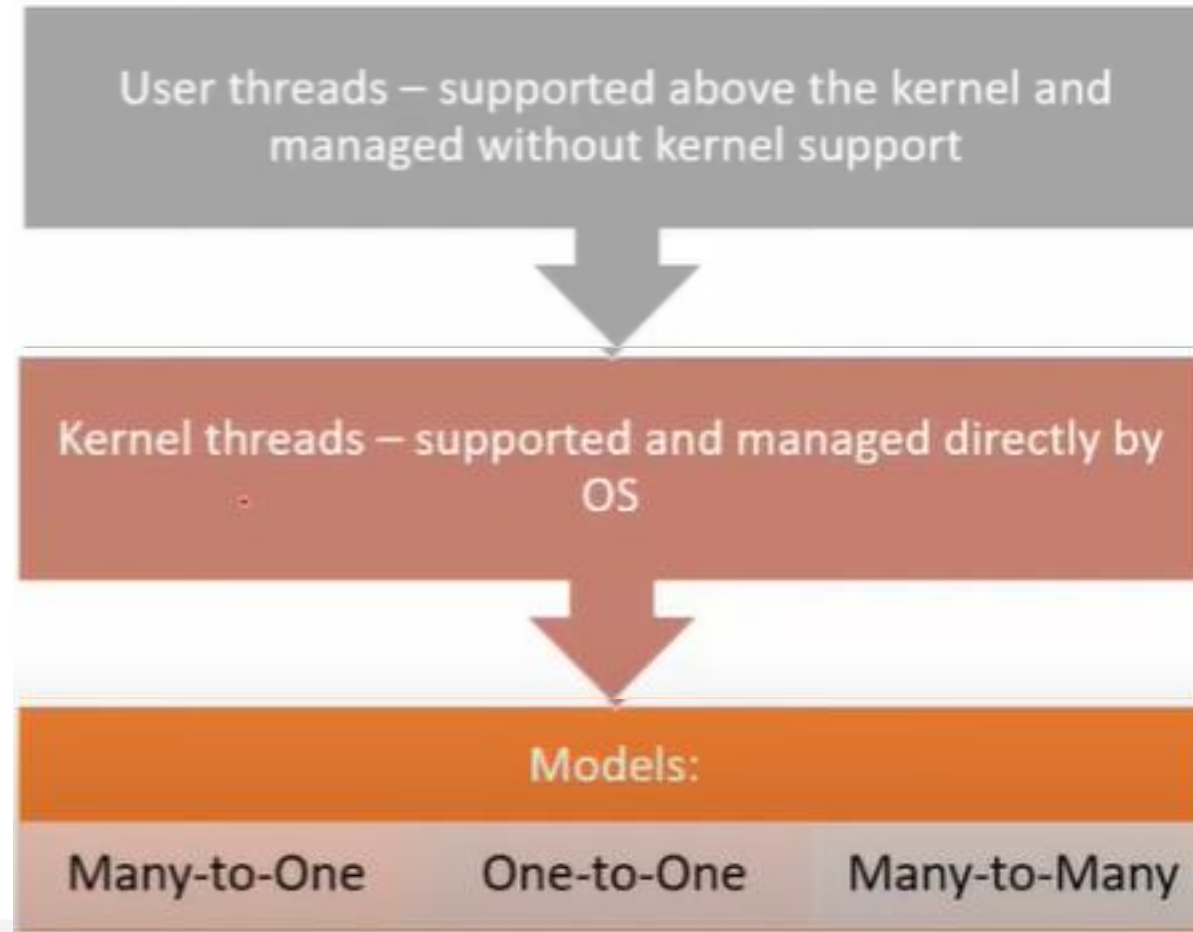


3. Multithreading Models

- Support for threads may be provided either at the user level, for **user threads**, or by the kernel, for **kernel threads**.
 - User threads are supported above the kernel and are managed without kernel support.
 - Kernel threads are supported and managed directly by the operating system.
 - Virtually all contemporary operating systems—including Windows, Linux, and macOS— support kernel threads.
 - Ultimately, a relationship must exist between user threads and kernel threads

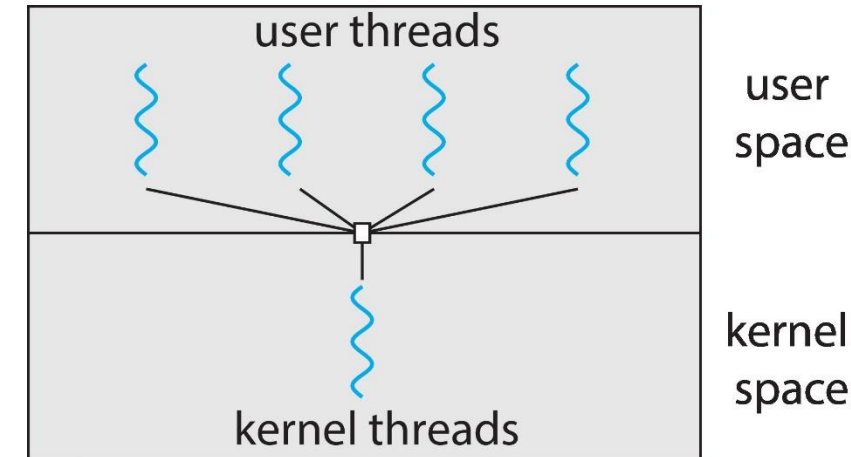


3. Multithreading Models



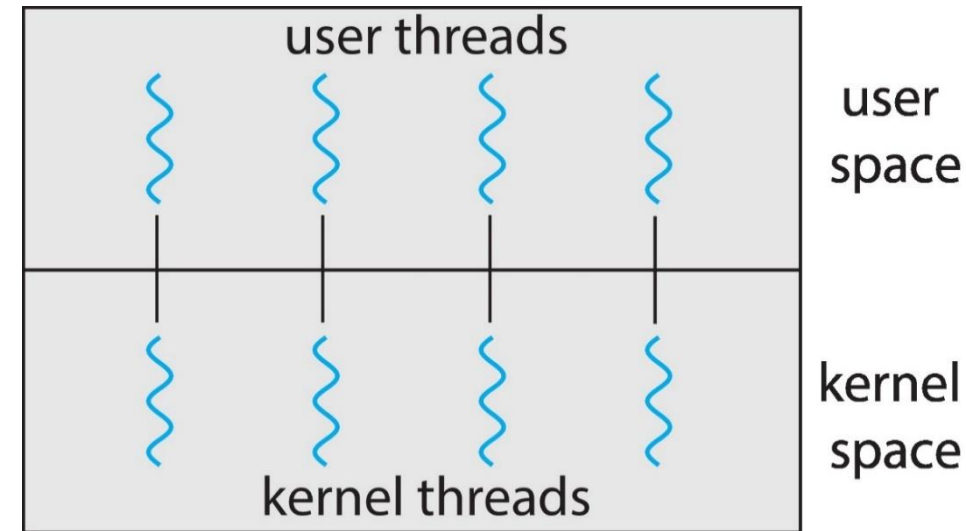
3. Multithreading Models

- Many-to-One Model
 - Many user-level threads mapped to single kernel thread.
 - Thread management is done by the thread library in user space, so it is efficient.
 - The entire process will block if a thread makes a blocking system call.
 - Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time.
 - Few systems currently use this model.
 - Solaris Green Threads, Since version 1.3, JVM is no longer implemented with green threads for any platform.
 - GNU Portable Thread



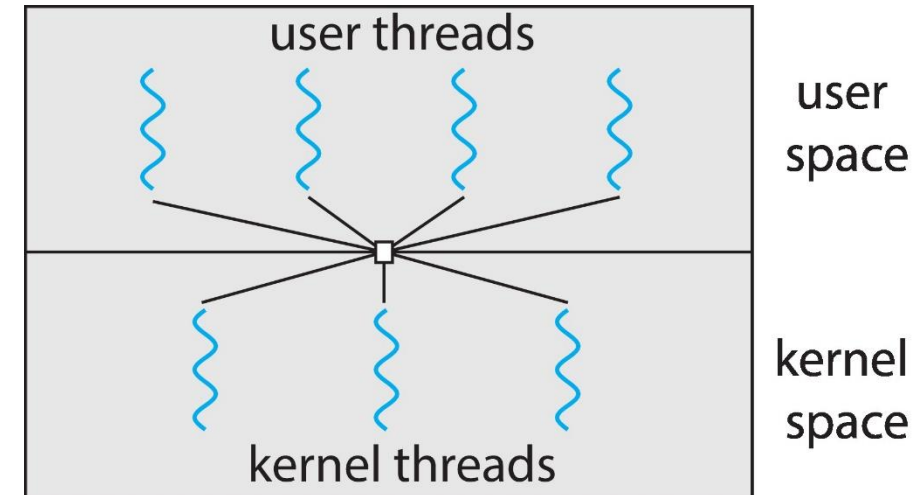
3. Multithreading Models

- One-to-One Model
 - Each user-level thread maps to kernel thread
 - Creating a user-level thread creates a kernel thread
 - More concurrency than many-to-one
 - Number of threads per process sometimes restricted due to overhead
 - Examples
 - Windows
 - Linux



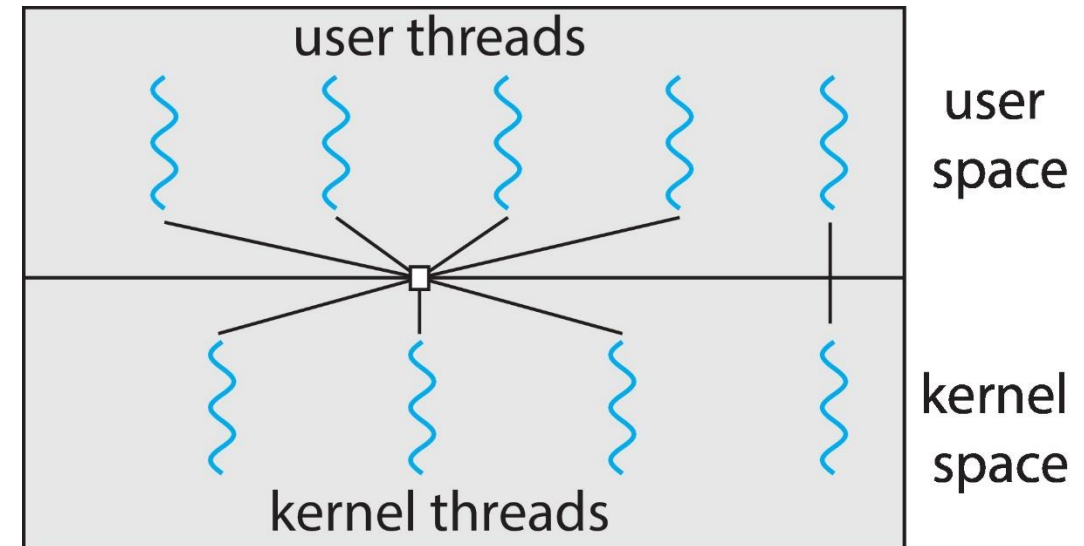
3. Multithreading Models

- Many-to-Many Model
 - Allows many user level threads to be mapped to many kernel threads
 - Allows the operating system to create a sufficient number of kernel threads
 - Windows with the ThreadFiber package
 - Otherwise not very common



3. Multithreading Models

- Two-level Model
 - Similar to Many-to-Many model, except that it allows a user thread to be bound to a kernel thread.
 - IRIX, HP-UX, and Tru64 UNIX use the two-tier model, as did Solaris prior to Solaris 9





4. Thread Libraries

- Thread library provides programmers with API for creating and managing threads.
- Two primary ways of implementing
 - Library entirely in user space
 - Code and data structures for the library exist in user space.
 - This means that invoking a function in the library results in a local function call in user space and not a system call.
 - Kernel-level library supported by the OS
 - Code and data structures for the library exist in kernel space.
 - Invoking a function in the API for the library typically results in a system call to the kernel.



4. Thread Libraries

- There are 3 main thread libraries in use today:
 - POSIX Pthreads - may be provided as either a user or kernel library, as an extension to the POSIX standard.
 - Win32 threads - provided as a kernel-level library on Windows systems.
 - Java threads - Since Java generally runs on a Java Virtual Machine, the implementation of threads is based upon whatever OS and hardware the JVM is running on, i.e. either Pthreads or Win32 threads depending on the system.



Pthreads

- Pthreads, the threads extension of the POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- May be provided as either a user-level or a kernel-level library
- Specification, not implementation
 - API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)



Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

- The program has two threads: the initial (or parent) thread in main() and the summation (or child) thread performing the summation operation in the runner() function.
- This program follows the thread create/join strategy, whereby after creating the summation thread, the parent thread will wait for it to terminate by calling the pthread join() function.
- The summation thread will terminate when it calls the function pthread exit(). Once the summation thread has returned, the parent thread will output the value of the shared data sum.



Pthreads Example

- Pthread code for joining ten threads.

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
 - On a UNIX system the JVM normally uses PThreads and on a Windows system it normally uses windows threads.
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```
- Standard practice is to implement Runnable interface



Java Threads

- Implementing Runnable interface:

```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

- Creating a thread:

```
Thread worker = new Thread(new Task());
worker.start();
```

- Waiting on a thread:

```
try {
    worker.join();
}
catch (InterruptedException ie) { }
```